

Ensemble

# Framework of Ensemble

- Get a set of classifiers

- $f_1(x), f_2(x), f_3(x), \dots$



坦



補



DD

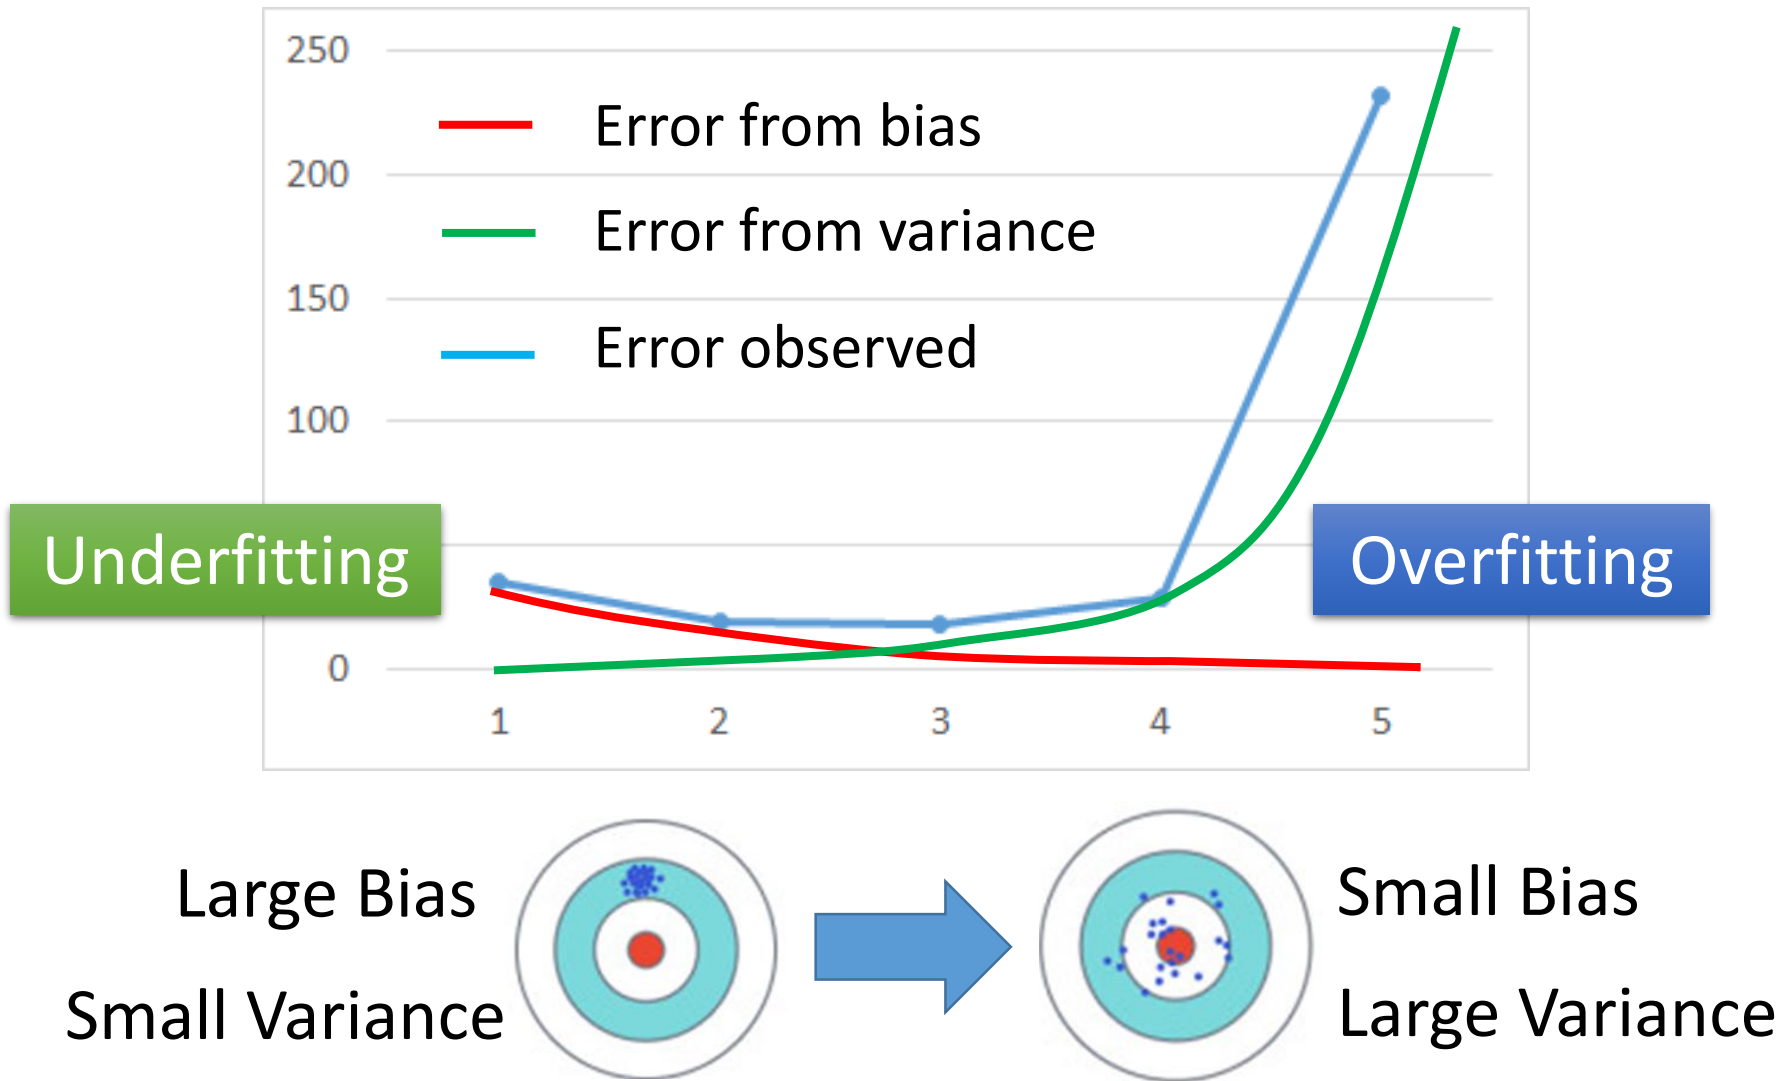
They should be diverse.

- Aggregate the classifiers (*properly*)

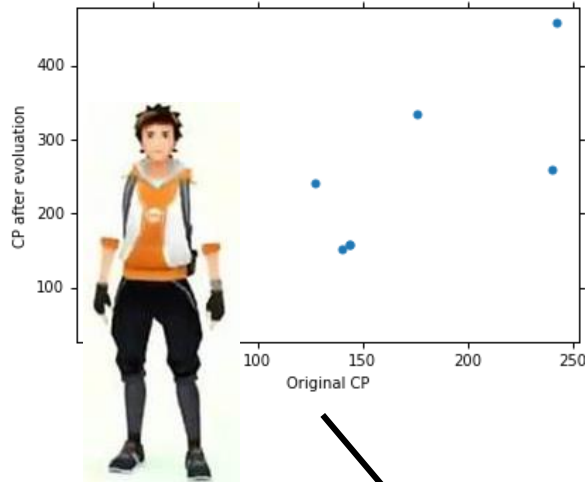
- 在打王時每個人都有該站的位置

# Ensemble: Bagging

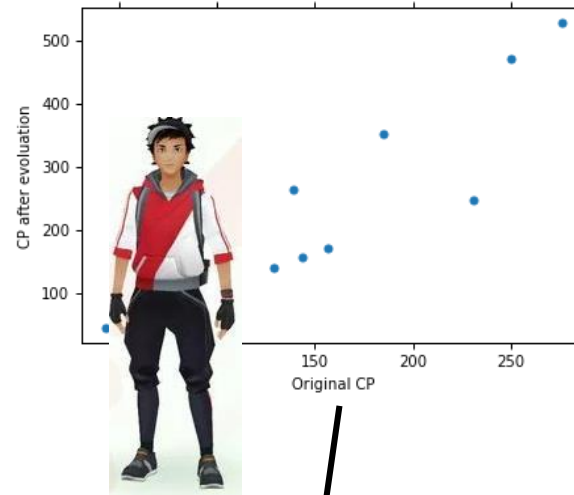
# Review: Bias v.s. Variance



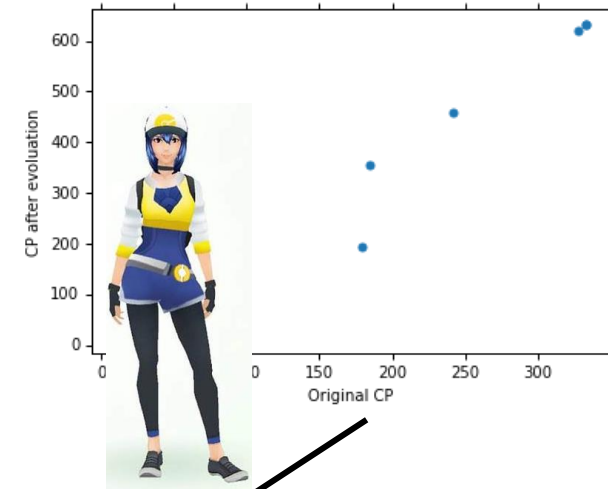
## Universe 1



## Universe 2

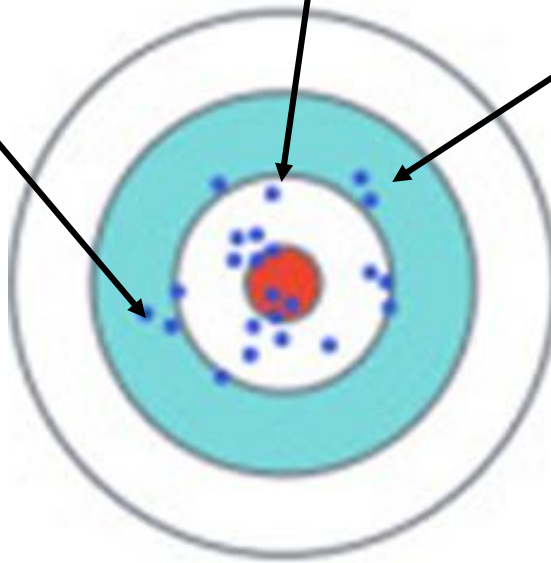


## Universe 3



A complex model will have large variance.

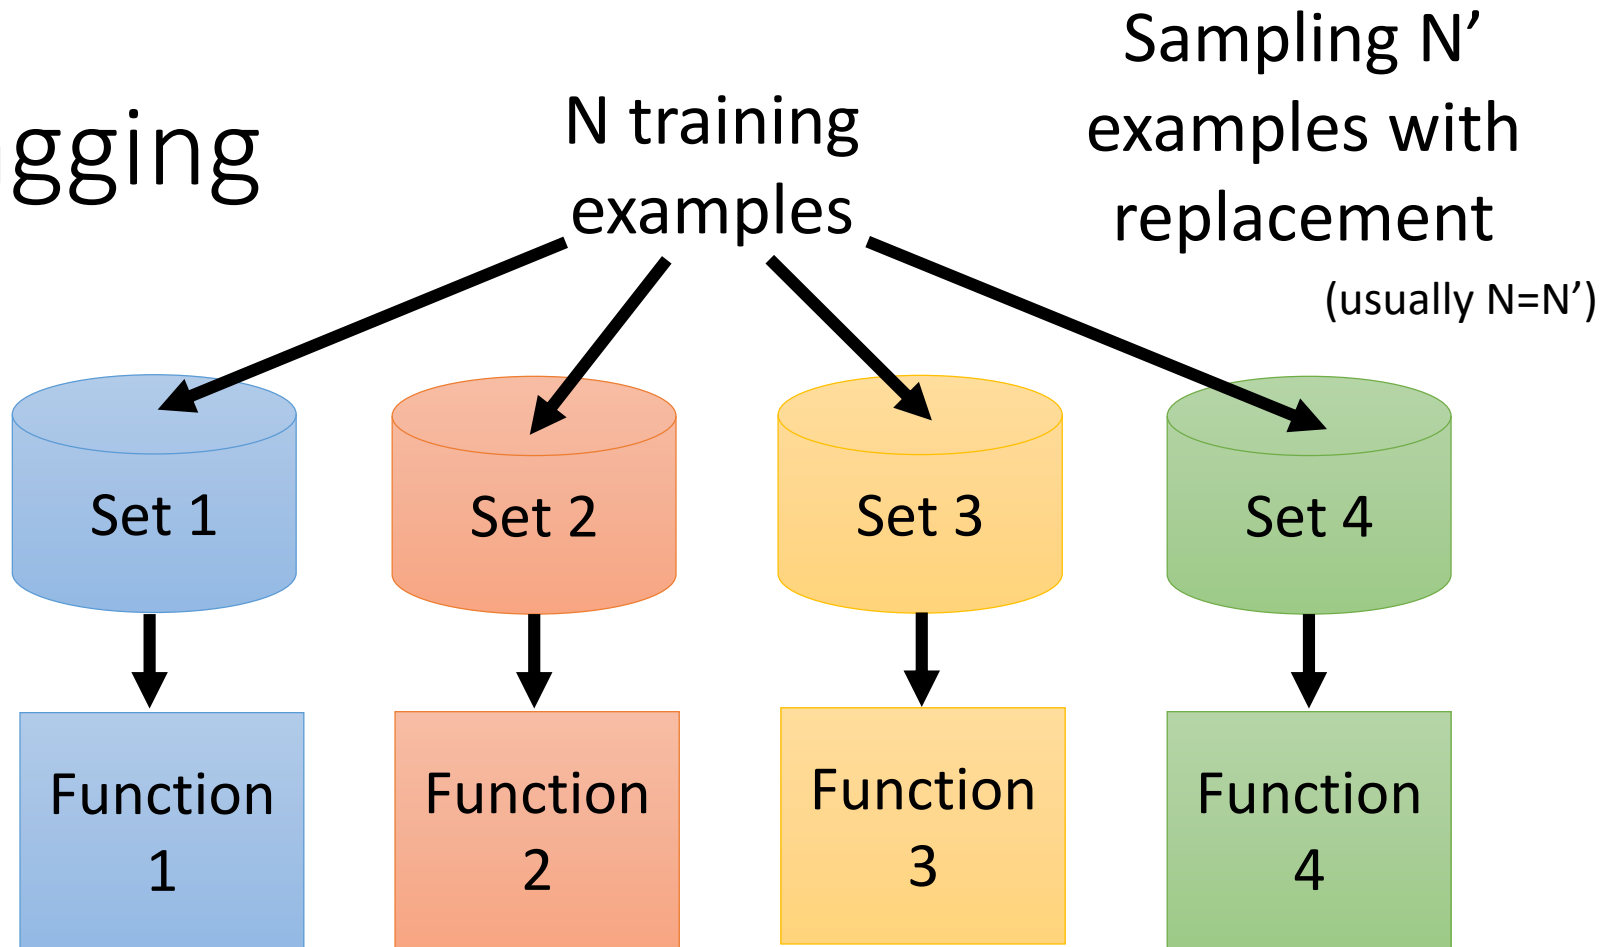
We can average complex models to reduce variance.



If we average all the  $f^*$ , is it close to  $\hat{f}$

$$E[f^*] = \hat{f}$$

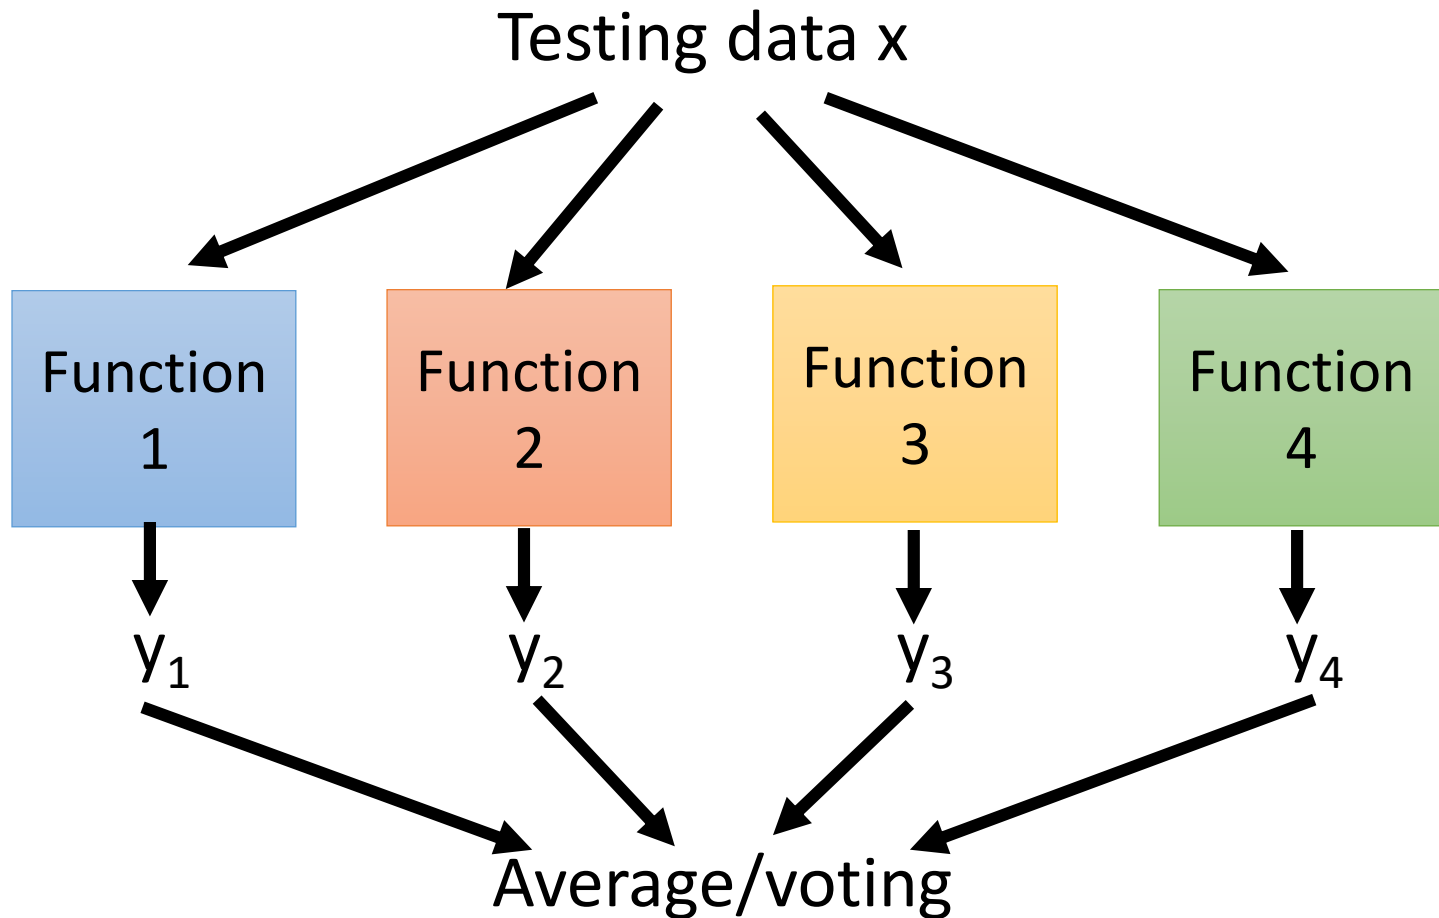
# Bagging



# Bagging

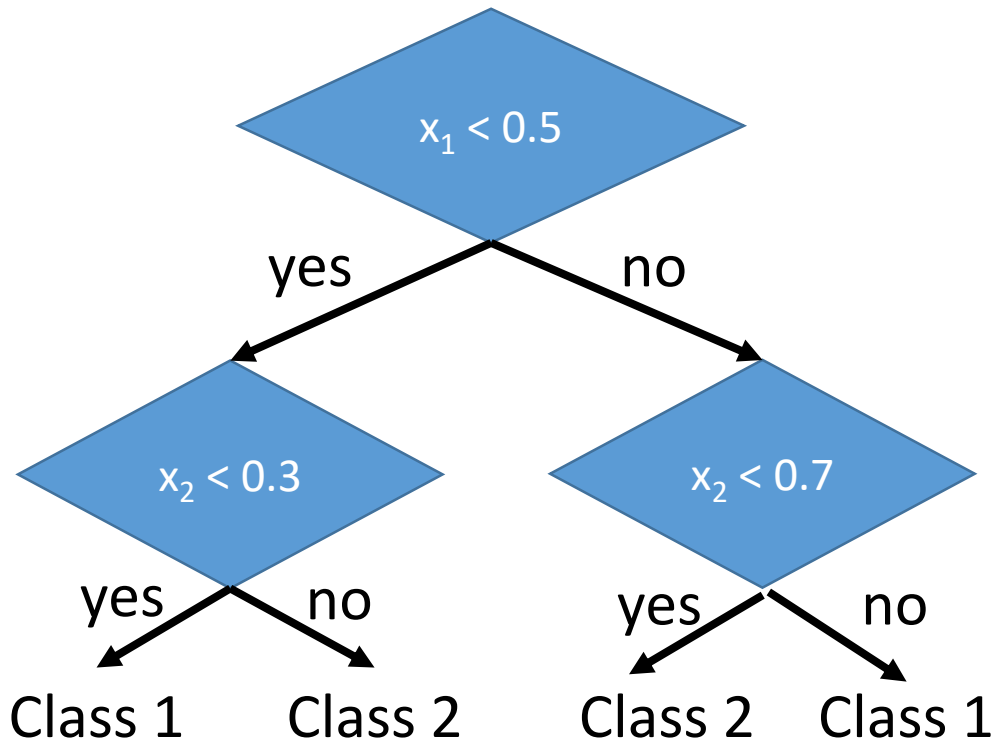
This approach would be helpful when  
your model is complex, easy to overfit.

e.g. decision tree

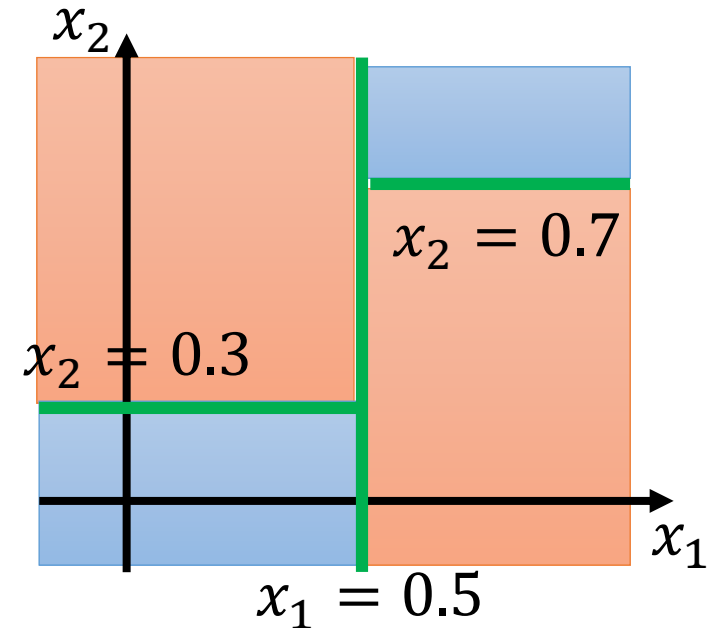


# Decision Tree

Assume each object  $x$  is represented by a 2-dim vector  $\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$



Can have more complex questions

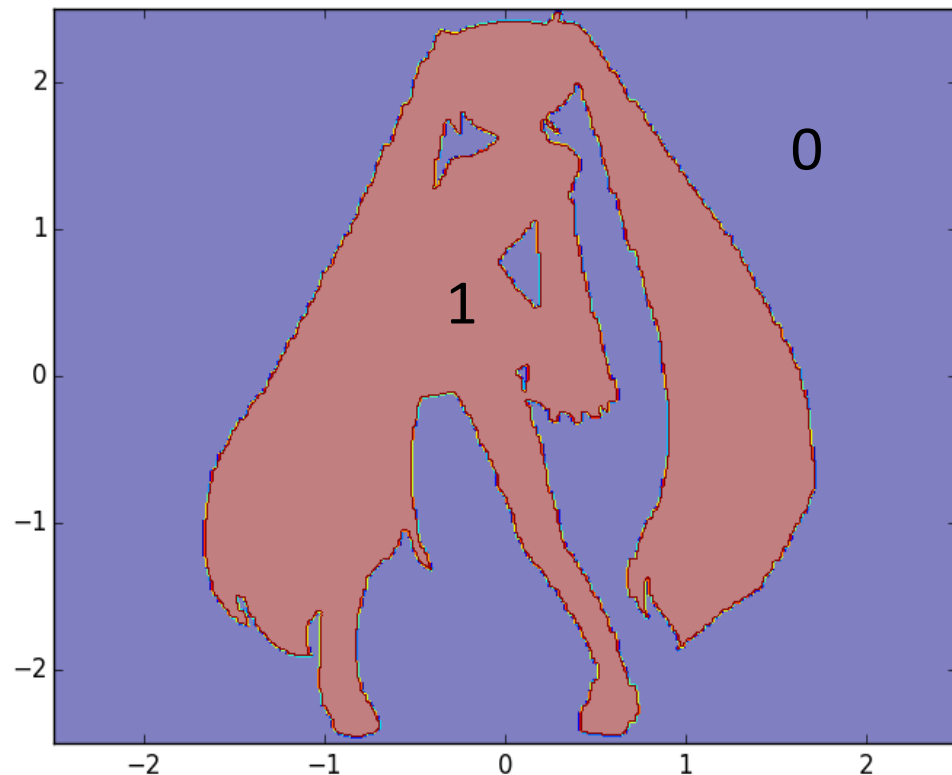


The questions in training .....

number of branches,  
Branching criteria,  
termination criteria,  
base hypothesis



# Experiment: Function of Miku

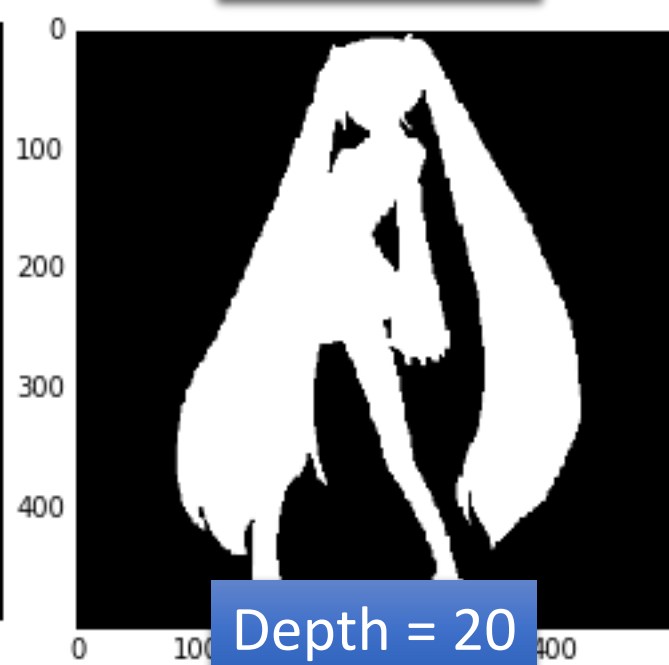
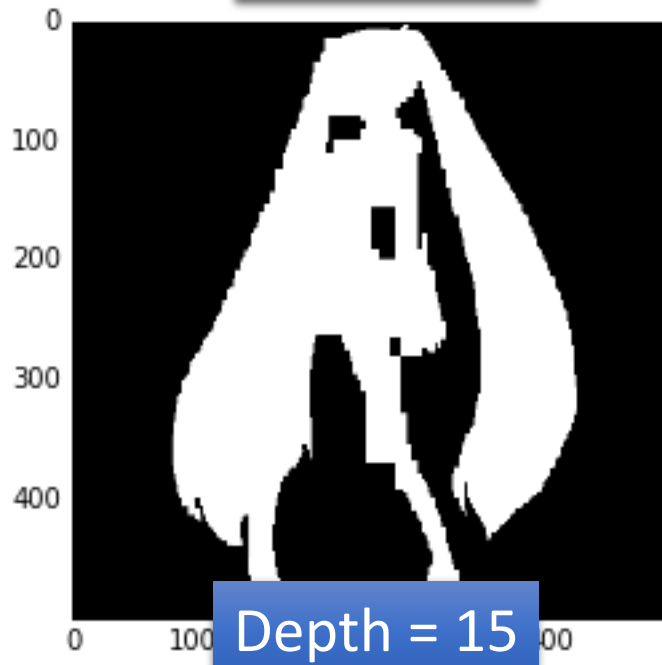
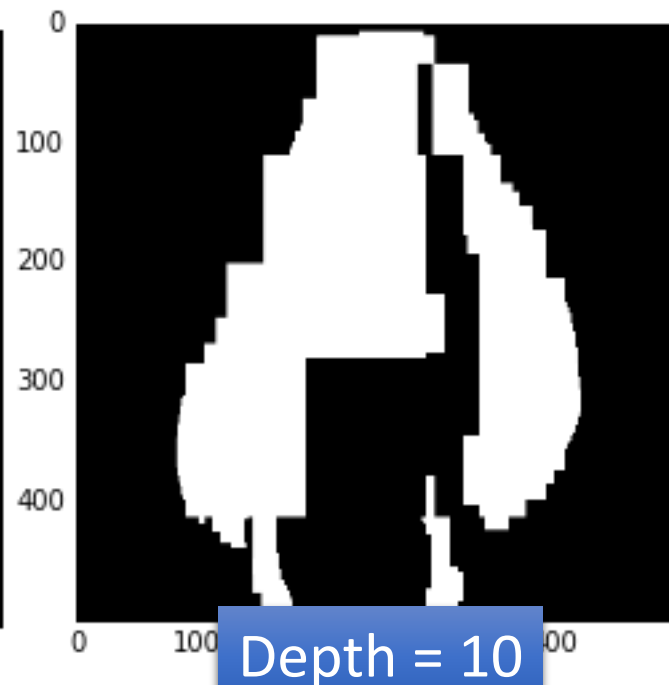
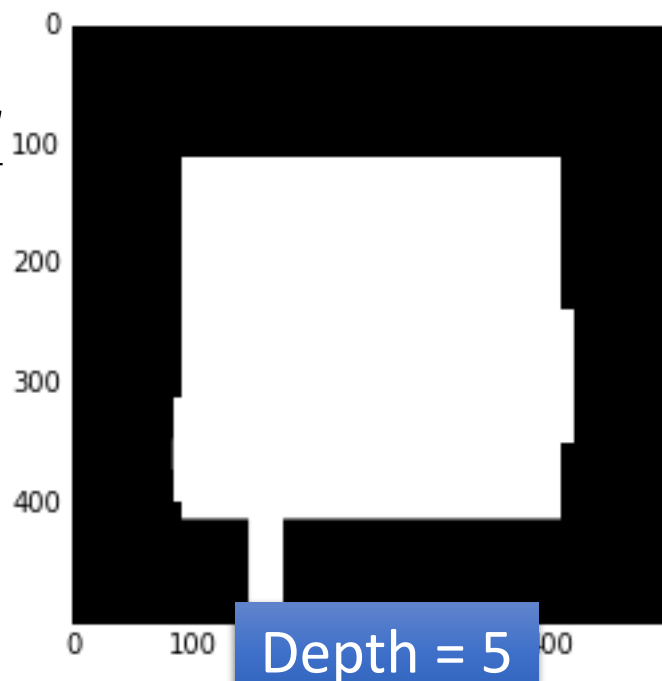


[http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS\\_2015\\_2/theano/miku](http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2/theano/miku)

(1<sup>st</sup> column: x, 2<sup>nd</sup> column: y, 3<sup>rd</sup> column: output (1 or 0) )

**Experiment:**  
**Function of Miku**

Single  
Decision  
Tree



# Random Forest

train	$f_1$	$f_2$	$f_3$	$f_4$
$x^1$	O	X	O	X
$x^2$	O	X	X	O
$x^3$	X	O	O	X
$x^4$	X	O	X	O

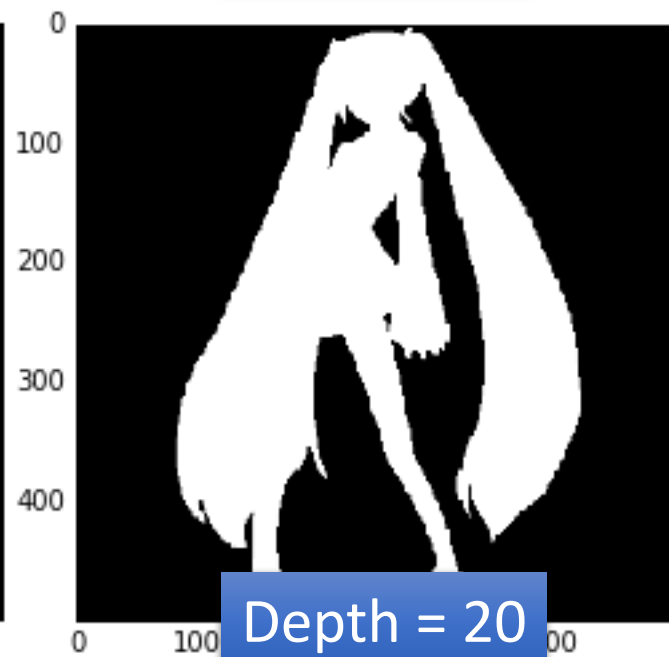
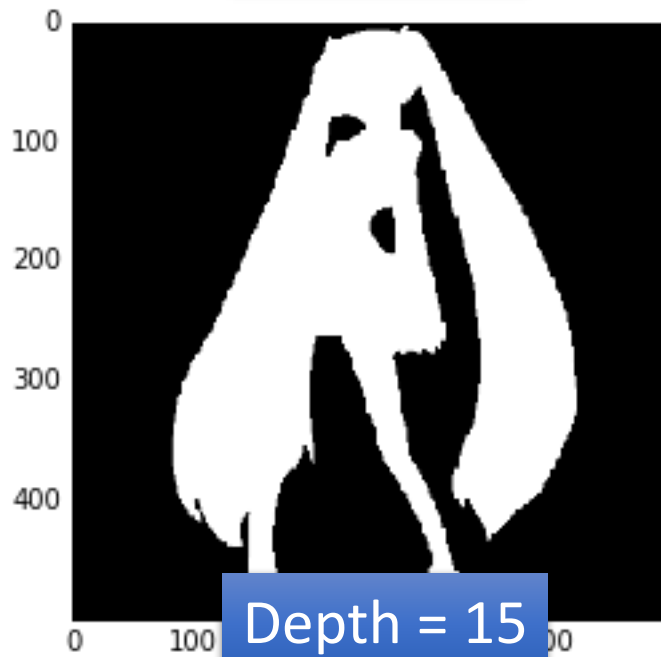
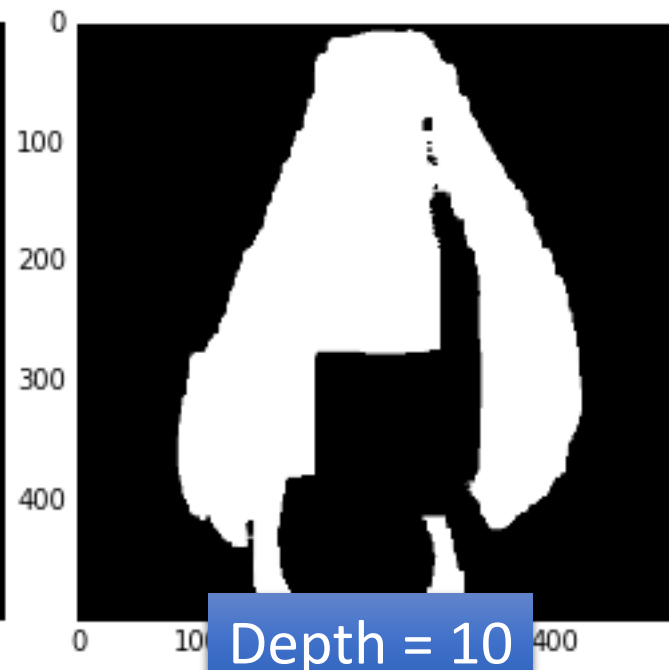
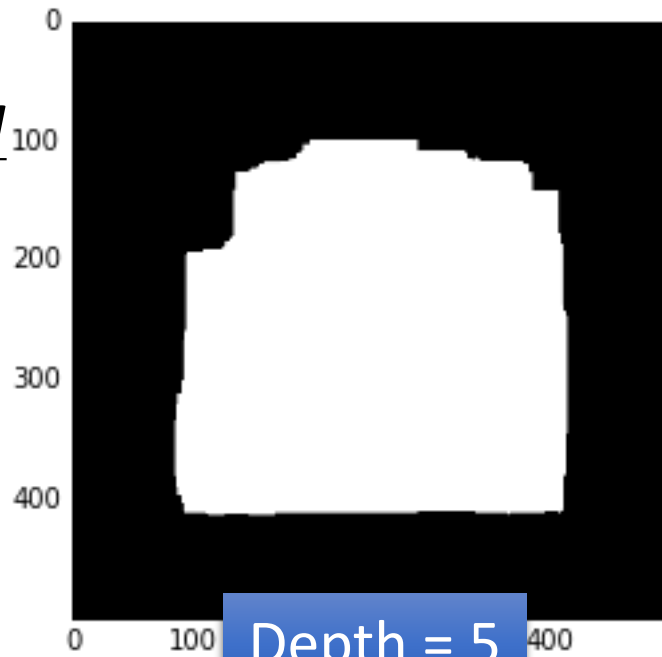
- Decision tree:
  - Easy to achieve 0% error rate on training data
    - If each training example has its own leaf .....
- Random forest: Bagging of decision tree
  - Resampling training data is not sufficient
  - Randomly restrict the features/questions used in each split
- Out-of-bag validation for bagging
  - Using RF =  $f_2 + f_4$  to test  $x^1$
  - Using RF =  $f_2 + f_3$  to test  $x^2$
  - Using RF =  $f_1 + f_4$  to test  $x^3$
  - Using RF =  $f_1 + f_3$  to test  $x^4$

Out-of-bag (OOB) error  
Good error estimation  
of testing set

**Experiment:**  
**Function of Miku**

Random  
Forest

(100 trees)



# Ensemble: Boosting

Improving Weak Classifiers

# Boosting

Training data:

$$\{(x^1, \hat{y}^1), \dots, (x^n, \hat{y}^n), \dots, (x^N, \hat{y}^N)\}$$

$\hat{y} = \pm 1$  (binary classification)

- Guarantee:
  - If your ML algorithm can produce classifier with error rate smaller than 50% on training data
  - You can obtain 0% error rate classifier after boosting.
- Framework of boosting
  - Obtain the first classifier  $f_1(x)$
  - Find another function  $f_2(x)$  to help  $f_1(x)$ 
    - However, if  $f_2(x)$  is similar to  $f_1(x)$ , it will not help a lot.
    - We want  $f_2(x)$  to be complementary with  $f_1(x)$  (How?)
  - Obtain the second classifier  $f_2(x)$
  - ..... Finally, combining all the classifiers
- The classifiers are learned sequentially.

# How to obtain different classifiers?

- Training on different training data sets
- How to have different training data sets
  - Re-sampling your training data to form a new set
  - Re-weighting your training data to form a new set
  - In real implementation, you only have to change the cost/objective function

$$(x^1, \hat{y}^1, u^1) \quad u^1 = \cancel{1} \quad 0.4$$

$$(x^2, \hat{y}^2, u^2) \quad u^2 = \cancel{1} \quad 2.1$$

$$(x^3, \hat{y}^3, u^3) \quad u^3 = \cancel{1} \quad 0.7$$

$$L(f) = \sum_n l(f(x^n), \hat{y}^n)$$



$$L(f) = \sum_n u^n l(f(x^n), \hat{y}^n)$$

# Idea of Adaboost

- Idea: **training  $f_2(x)$  on the new training set that fails  $f_1(x)$**
- How to find a new training set that fails  $f_1(x)$ ?

$\varepsilon_1$ : the error rate of  $f_1(x)$  on its training data

$$\varepsilon_1 = \frac{\sum_n u_1^n \delta(f_1(x^n) \neq \hat{y}^n)}{Z_1} \quad Z_1 = \sum_n u_1^n \quad \varepsilon_1 < 0.5$$

Changing the example weights from  $u_1^n$  to  $u_2^n$  such that

$$\frac{\sum_n u_2^n \delta(f_1(x^n) \neq \hat{y}^n)}{Z_2} = 0.5$$

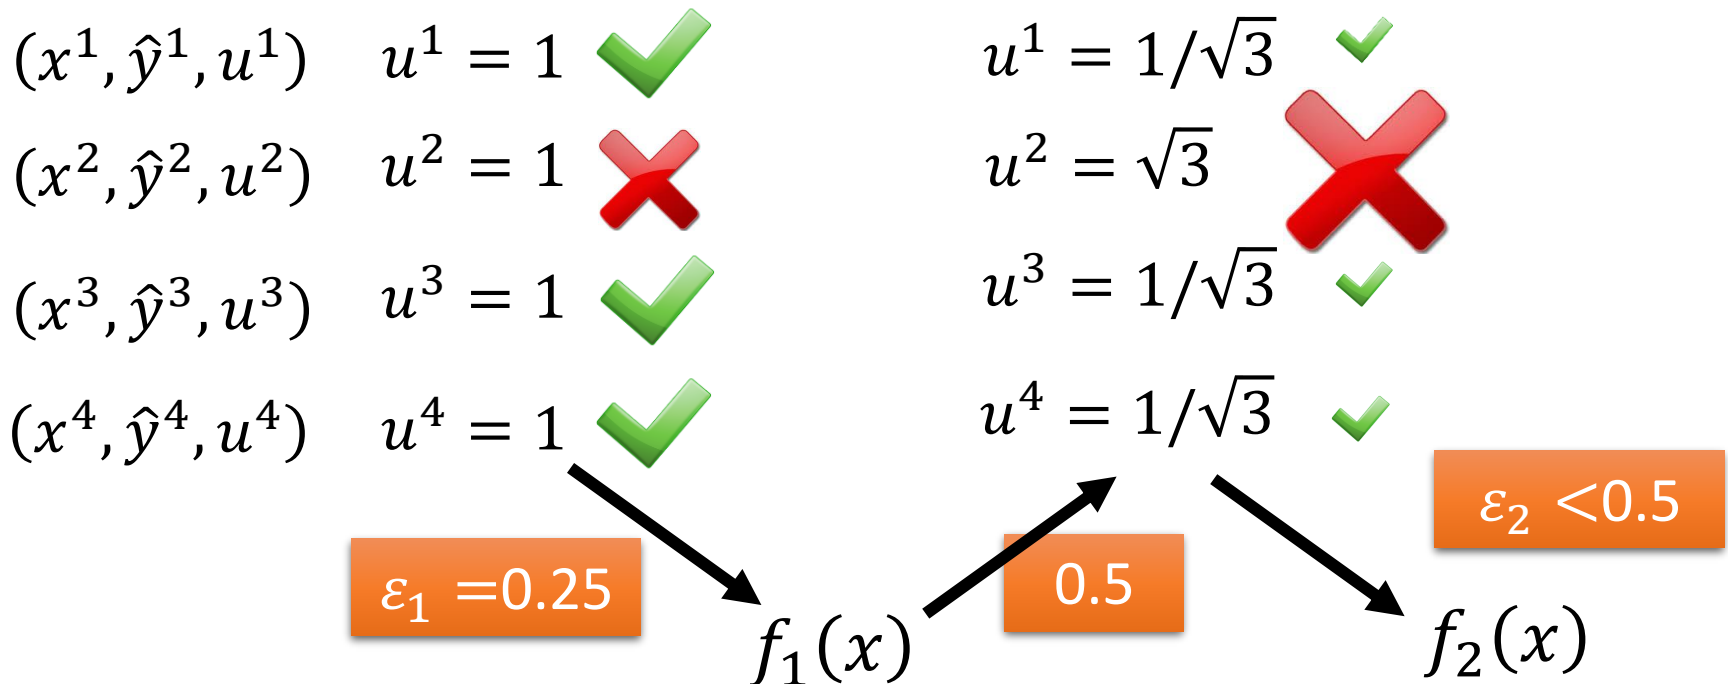
The performance of  $f_1$  for new weights would be random.

Training  $f_2(x)$  based on the new weights  $u_2^n$



# Re-weighting Training Data

- Idea: training  $f_2(x)$  on the new training set that fails  $f_1(x)$
- How to find a new training set that fails  $f_1(x)$ ?



# Re-weighting Training Data

- Idea: **training  $f_2(x)$  on the new training set that fails  $f_1(x)$**
- How to find a new training set that fails  $f_1(x)$ ?

$$\left\{ \begin{array}{ll} \text{If } x^n \text{ misclassified by } f_1 \ (f_1(x^n) \neq \hat{y}^n) & u_2^n \leftarrow u_1^n \text{ multiplying } d_1 \quad \text{increase} \\ \text{If } x^n \text{ correctly classified by } f_1 \ (f_1(x^n) = \hat{y}^n) & u_2^n \leftarrow u_1^n \text{ divided by } d_1 \quad \text{decrease} \end{array} \right.$$

$f_2$  will be learned based on example weights  $u_2^n$

What is the value of  $d_1$ ?

# Re-weighting Training Data

$$\varepsilon_1 = \frac{\sum_n u_1^n \delta(f_1(x^n) \neq \hat{y}^n)}{Z_1}$$

$$Z_1 = \sum_n u_1^n$$

$$\frac{\sum_n u_2^n \delta(f_1(x^n) \neq \hat{y}^n)}{Z_2} = 0.5$$

$$\begin{aligned} f_1(x^n) \neq \hat{y}^n & \quad u_2^n \leftarrow u_1^n \text{ multiplying } d_1 \\ f_1(x^n) = \hat{y}^n & \quad u_2^n \leftarrow u_1^n \text{ divided by } d_1 \end{aligned}$$

$$= \sum_{f_1(x^n) \neq \hat{y}^n} u_1^n d_1$$

$$= \sum_{f_1(x^n) \neq \hat{y}^n} u_2^n + \sum_{f_1(x^n) = \hat{y}^n} u_2^n$$

$$= \sum_n u_2^n$$

$$= \sum_{f_1(x^n) \neq \hat{y}^n} u_1^n d_1 + \sum_{f_1(x^n) = \hat{y}^n} u_1^n / d_1$$

$$\frac{\sum_{f_1(x^n) \neq \hat{y}^n} u_1^n d_1 + \sum_{f_1(x^n) = \hat{y}^n} u_1^n / d_1}{\sum_{f_1(x^n) \neq \hat{y}^n} u_1^n d_1} = 2$$

# Re-weighting Training Data

$$\varepsilon_1 = \frac{\sum_n u_1^n \delta(f_1(x^n) \neq \hat{y}^n)}{Z_1}$$

$$Z_1 = \sum_n u_1^n$$

$$\frac{\sum_n u_2^n \delta(f_1(x^n) \neq \hat{y}^n)}{Z_2} = 0.5$$

$$\begin{aligned} f_1(x^n) \neq \hat{y}^n & \quad u_2^n \leftarrow u_1^n \text{ multiplying } d_1 \\ f_1(x^n) = \hat{y}^n & \quad u_2^n \leftarrow u_1^n \text{ divided by } d_1 \end{aligned}$$

$$\frac{\sum_{f_1(x^n) \neq \hat{y}^n} u_1^n d_1 + \sum_{f_1(x^n) = \hat{y}^n} u_1^n / d_1}{\sum_{f_1(x^n) \neq \hat{y}^n} u_1^n d_1} = 2 \quad \frac{\sum_{f_1(x^n) = \hat{y}^n} u_1^n / d_1}{\sum_{f_1(x^n) \neq \hat{y}^n} u_1^n d_1} = 1$$

$$\sum_{f_1(x^n) = \hat{y}^n} u_1^n / d_1 = \sum_{f_1(x^n) \neq \hat{y}^n} u_1^n d_1 \quad \frac{1}{d_1} \sum_{f_1(x^n) = \hat{y}^n} u_1^n = d_1 \sum_{f_1(x^n) \neq \hat{y}^n} u_1^n$$

$$\varepsilon_1 = \frac{\sum_{f_1(x^n) \neq \hat{y}^n} u_1^n}{Z_1} \quad \frac{Z_1(1 - \varepsilon_1)}{Z_1 \varepsilon_1}$$

$$\sum_{f_1(x^n) \neq \hat{y}^n} u_1^n = Z_1 \varepsilon_1$$

$$Z_1(1 - \varepsilon_1)/d_1 = Z_1 \varepsilon_1 d_1$$

$$d_1 = \sqrt{(1 - \varepsilon_1)/\varepsilon_1} > 1$$

# Algorithm for AdaBoost

- Giving training data  $\{(x^1, \hat{y}^1, u_1^1), \dots, (x^n, \hat{y}^n, u_1^n), \dots, (x^N, \hat{y}^N, u_1^N)\}$ 
    - $\hat{y} = \pm 1$  (Binary classification),  $u_1^n = 1$  (equal weights)
  - For  $t = 1, \dots, T$ :
    - Training weak classifier  $f_t(x)$  with weights  $\{u_t^1, \dots, u_t^N\}$
    - $\varepsilon_t$  is the error rate of  $f_t(x)$  with weights  $\{u_t^1, \dots, u_t^N\}$
    - For  $n = 1, \dots, N$ :
      - If  $x^n$  is misclassified by  $f_t(x)$ :  $\hat{y}^n \neq f_t(x^n)$
      - $u_{t+1}^n = u_t^n \times d_t = u_t^n \times \exp(\alpha_t)$   $d_t = \sqrt{(1 - \varepsilon_t)/\varepsilon_t}$
      - Else:
      - $u_{t+1}^n = u_t^n / d_t = u_t^n \times \exp(-\alpha_t)$   $\alpha_t = \ln \sqrt{(1 - \varepsilon_t)/\varepsilon_t}$
- $$u_{t+1}^n \leftarrow u_t^n \times \exp(-\alpha_t)$$

# Algorithm for AdaBoost

- We obtain a set of functions:  $f_1(x), \dots, f_t(x), \dots, f_T(x)$
- How to aggregate them?
  - Uniform weight:
    - $H(x) = \text{sign}(\sum_{t=1}^T f_t(x))$
  - Non-uniform weight:
    - $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t f_t(x))$

Smaller error  $\varepsilon_t$ ,  
larger weight for  
final voting

$$\alpha_t = \ln \sqrt{(1 - \varepsilon_t) / \varepsilon_t}$$

$$\varepsilon_t = 0.1$$

$$\varepsilon_t = 0.4$$

$$u_{t+1}^n = u_t^n \times \exp(-\hat{y}^n f_t(x^n) \alpha_t)$$

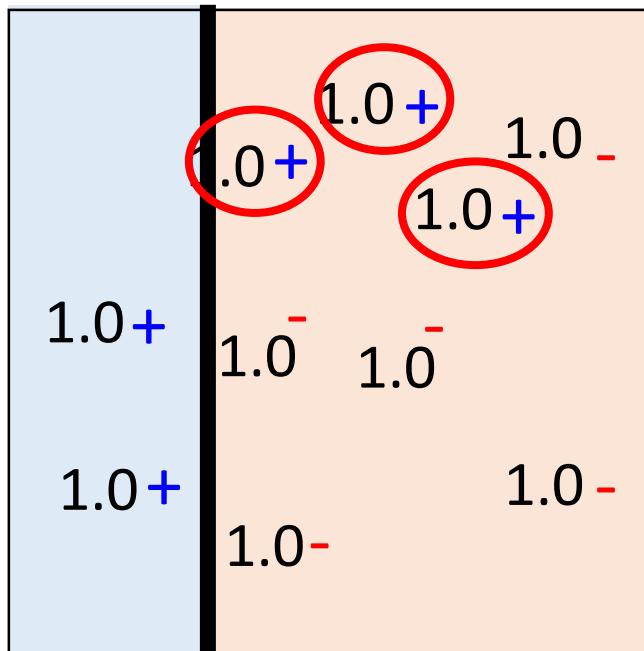
$$\alpha_t = 1.10$$

$$\alpha_t = 0.20$$

# Toy Example

$T=3$ , weak classifier = decision stump

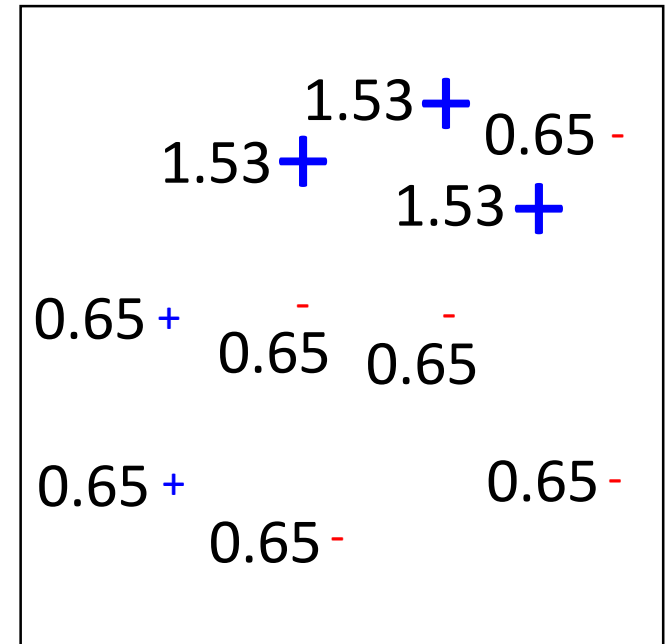
- $t=1$



$$\varepsilon_1 = 0.30$$

$$d_1 = 1.53$$

$$\alpha_1 = 0.42$$



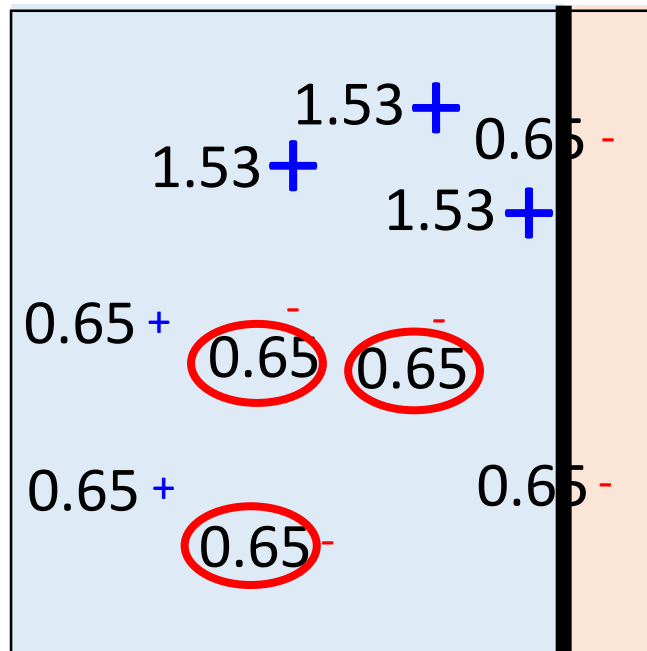
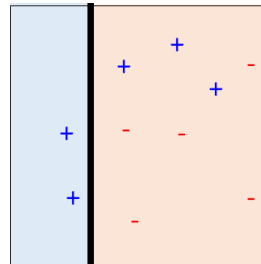
# Toy Example

$T=3$ , weak classifier = decision stump

•  $t=2$

$$f_1(x):$$

$$\alpha_1 = 0.42$$



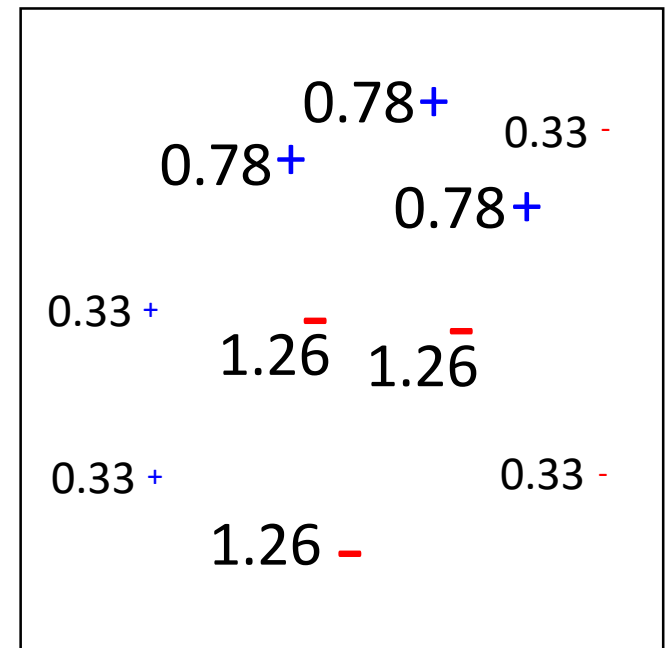
$$f_2(x)$$



$$\epsilon_2 = 0.21$$

$$d_2 = 1.94$$

$$\alpha_2 = 0.66$$





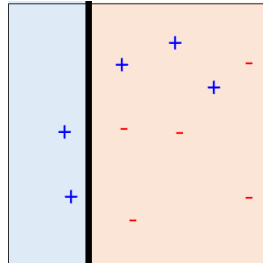
# Toy Example

$T=3$ , weak classifier = decision stump

•  $t=3$

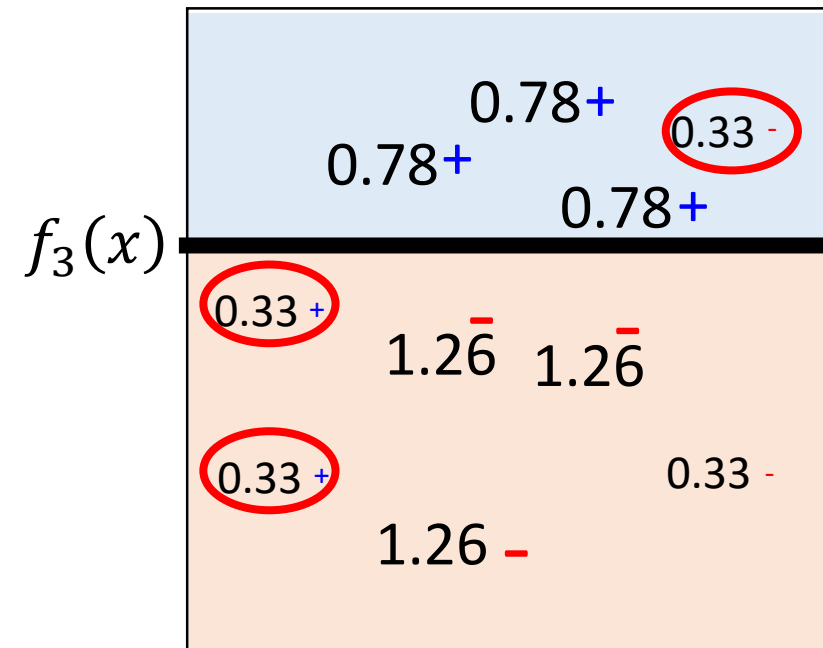
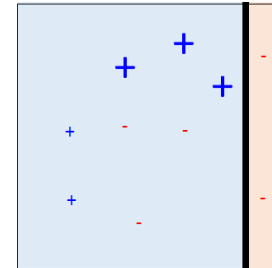
$f_1(x)$ :

$$\alpha_1 = 0.42$$



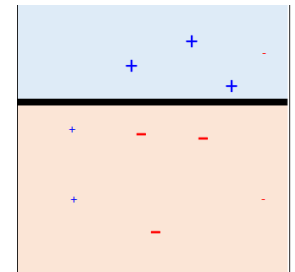
$f_2(x)$ :

$$\alpha_2 = 0.66$$



$f_3(x)$ :

$$\alpha_3 = 0.95$$



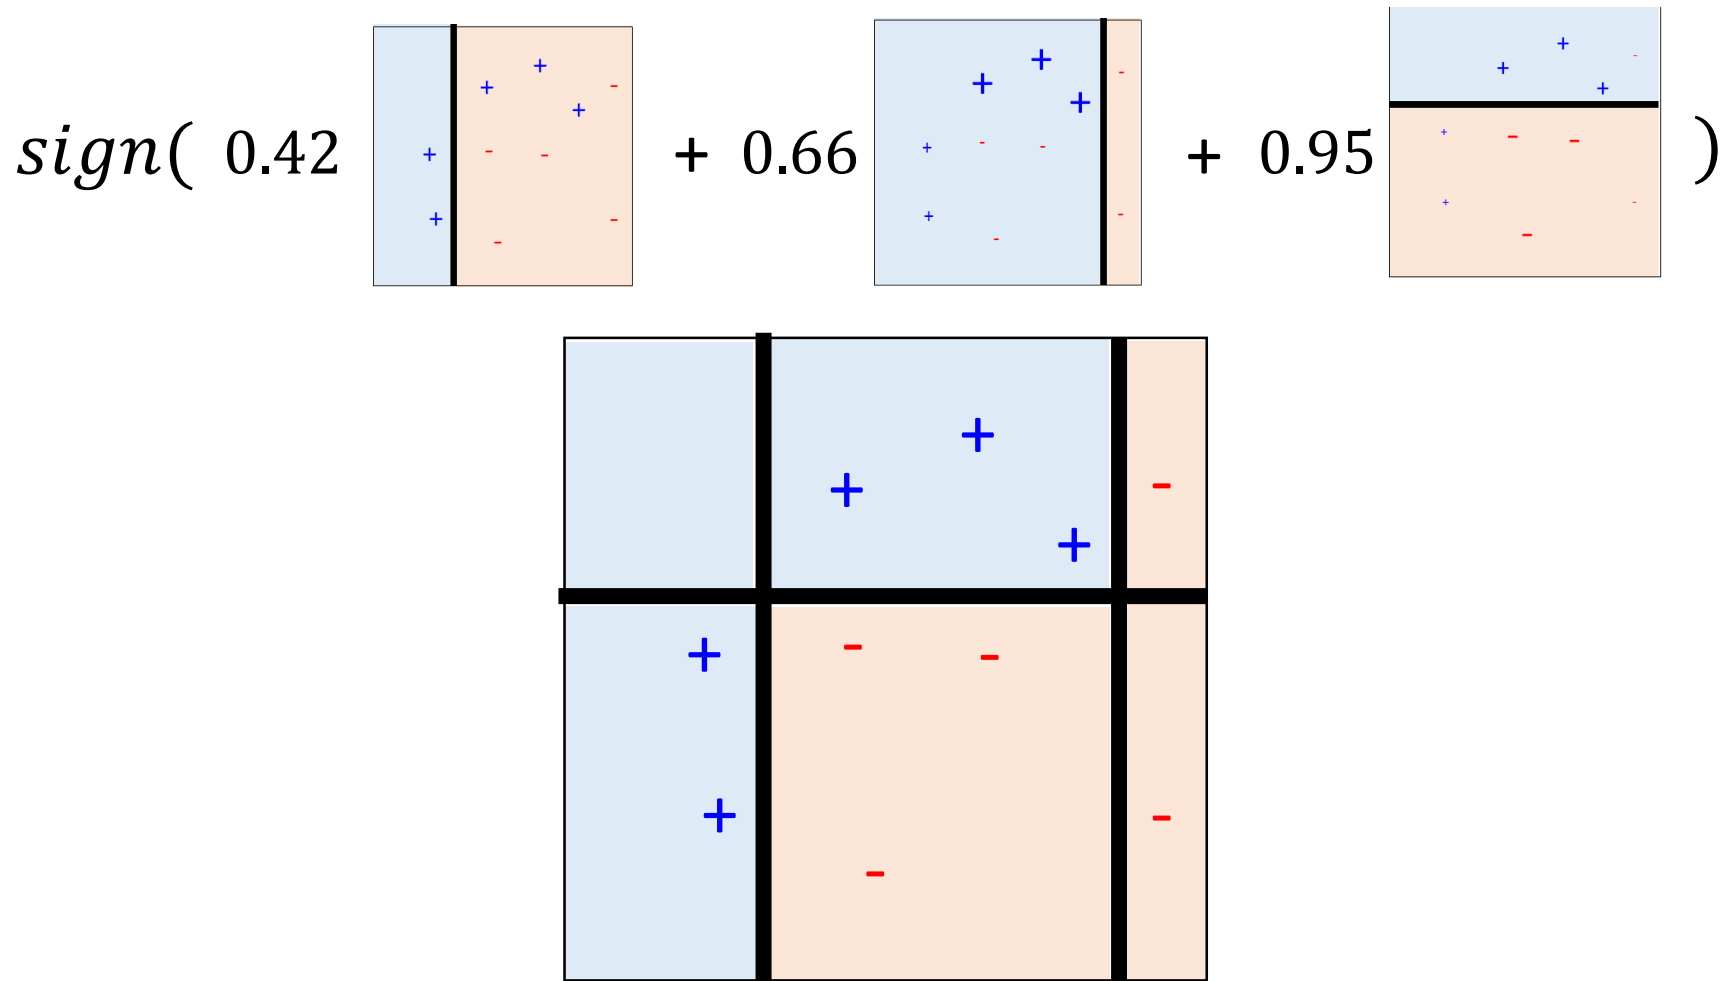
$$\varepsilon_3 = 0.13$$

$$d_3 = 2.59$$

$$\alpha_3 = 0.95$$

# Toy Example

- Final Classifier:  $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t f_t(x))$



# Warning of Math

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t f_t(x) \right) \quad \alpha_t = \ln \sqrt{(1 - \varepsilon_t) / \varepsilon_t}$$

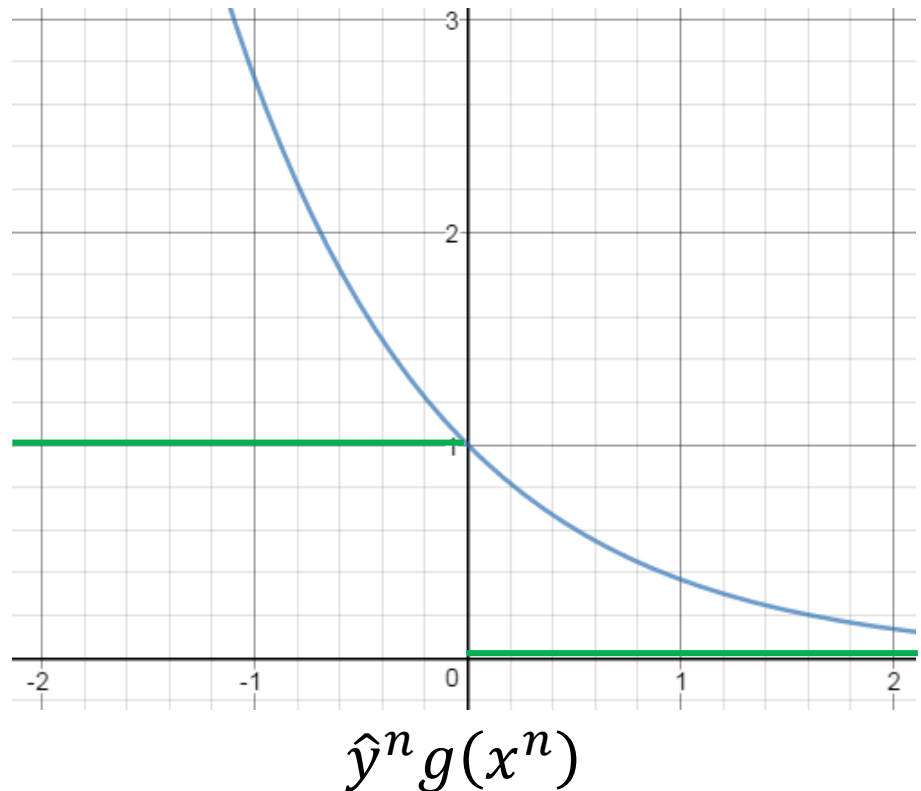
As we have more and more  $f_t$  (T increases),  $H(x)$  achieves smaller and smaller error rate on training data.

# Error Rate of Final Classifier

- Final classifier:  $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t f_t(x))$ 
  - $\alpha_t = \ln \sqrt{(1 - \varepsilon_t) / \varepsilon_t}$

Training Data Error Rate

$$\begin{aligned} &= \frac{1}{N} \sum_n \delta(H(x^n) \neq \hat{y}^n) \\ &= \frac{1}{N} \sum_n \delta(\hat{y}^n g(x^n) < 0) \\ &\leq \frac{1}{N} \sum_n \exp(-\hat{y}^n g(x^n)) \end{aligned}$$



Training Data Error Rate

$$\leq \frac{1}{N} \sum_n \exp(-\hat{y}^n g(x^n)) = \frac{1}{N} Z_{T+1}$$

$$g(x) = \sum_{t=1}^T \alpha_t f_t(x)$$

$$\alpha_t = \ln \sqrt{(1 - \varepsilon_t) / \varepsilon_t}$$

$Z_t$ : the summation of the weights of training data for training  $f_t$

What is  $Z_{T+1}$  =?  $Z_{T+1} = \sum_n u_{T+1}^n$

$$\left. \begin{array}{l} u_1^n = 1 \\ u_{t+1}^n = u_t^n \times \exp(-\hat{y}^n f_t(x^n) \alpha_t) \end{array} \right\} u_{T+1}^n = \prod_{t=1}^T \exp(-\hat{y}^n f_t(x^n) \alpha_t)$$

$$Z_{T+1} = \sum_n \prod_{t=1}^T \exp(-\hat{y}^n f_t(x^n) \alpha_t)$$

$$= \sum_n \exp \left( -\hat{y}^n \sum_{t=1}^T f_t(x^n) \alpha_t \right)$$

Training Data Error Rate

$$\leq \frac{1}{N} \sum_n \exp(-\hat{y}^n g(x^n)) = \frac{1}{N} Z_{T+1}$$

$$g(x) = \sum_{t=1}^T \alpha_t f_t(x)$$
$$\alpha_t = \ln \sqrt{(1 - \epsilon_t) / \epsilon_t}$$

$$Z_1 = N \quad (\text{equal weights})$$

$$Z_t = \underline{Z_{t-1} \epsilon_t} \exp(\alpha_t) + \underline{Z_{t-1} (1 - \epsilon_t)} \exp(-\alpha_t)$$

Misclassified portion in  $Z_{t-1}$

Correctly classified portion in  $Z_{t-1}$

$$= Z_{t-1} \epsilon_t \sqrt{(1 - \epsilon_t) / \epsilon_t} + Z_{t-1} (1 - \epsilon_t) \sqrt{\epsilon_t / (1 - \epsilon_t)}$$

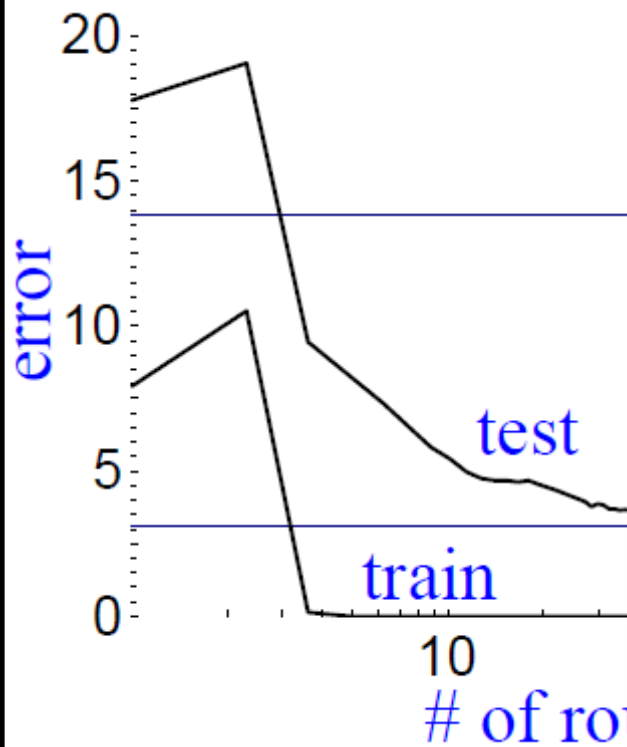
$$= Z_{t-1} \times 2\sqrt{\epsilon_t (1 - \epsilon_t)} \quad Z_{T+1} = N \prod_{t=1}^T 2\sqrt{\epsilon_t (1 - \epsilon_t)}$$

$$\text{Training Data Error Rate} \leq \prod_{t=1}^T \underline{2\sqrt{\epsilon_t (1 - \epsilon_t)}}_{<1}$$

Smaller and  
smaller

End of Warning

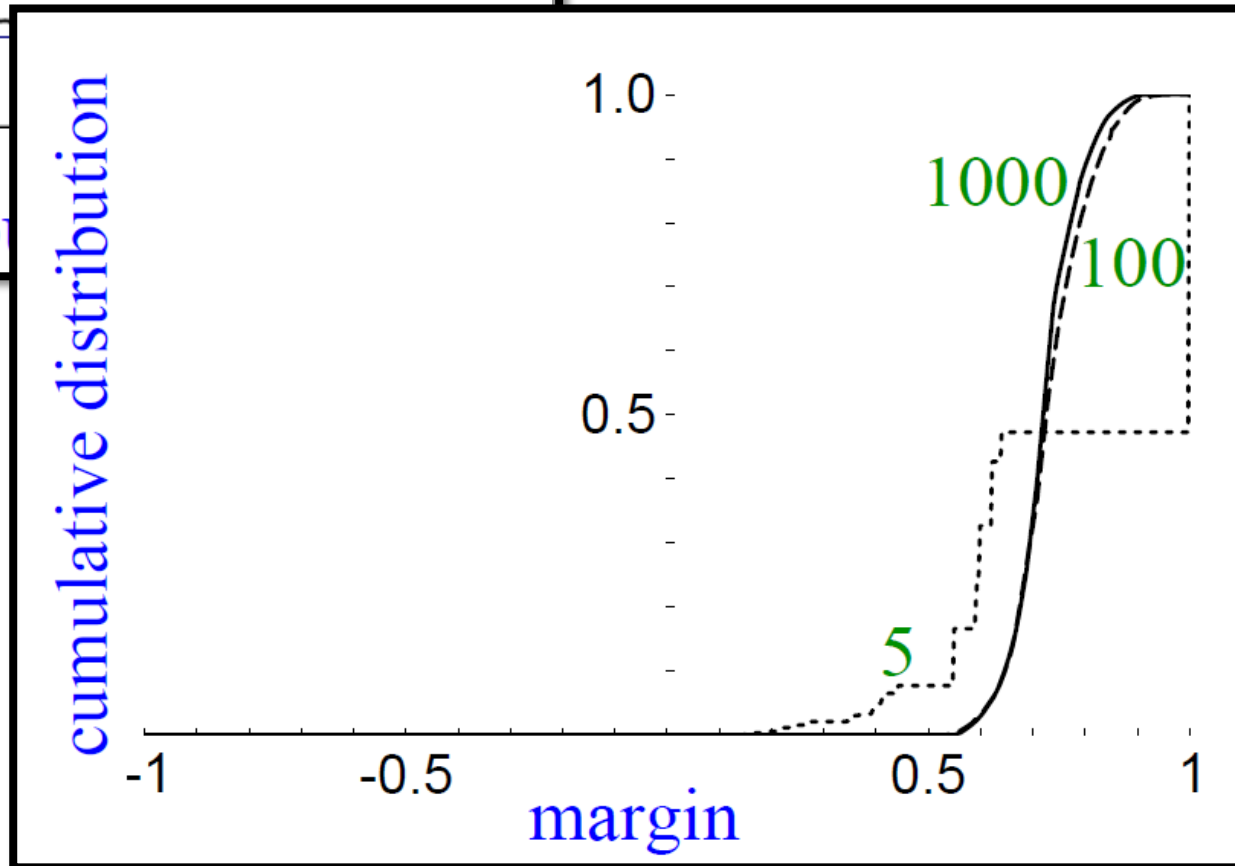
Even though the training error is 0, the testing error still decreases?



$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t f_t(x) \right)$$

$g(x)$

$$\text{Margin} = \hat{y} g(x)$$





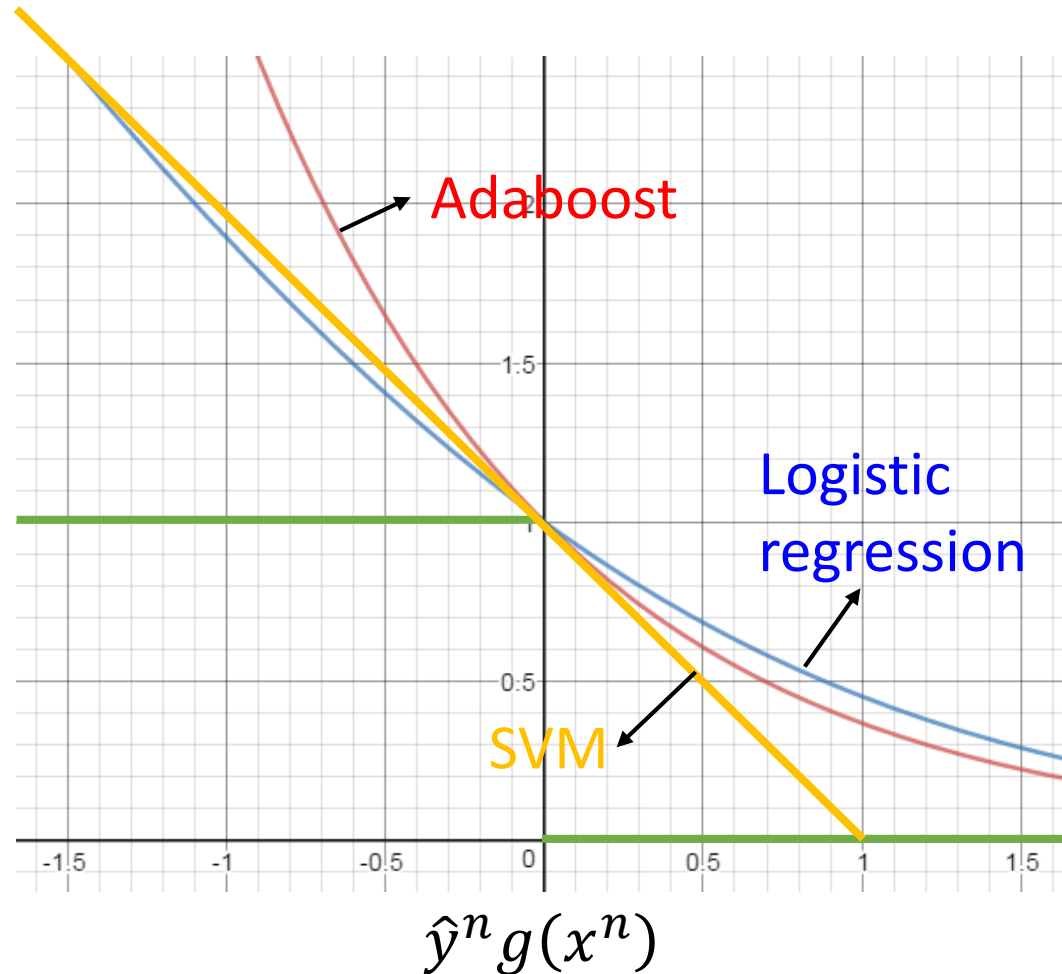
# Large Margin?

$$H(x) = \text{sign} \left( \underbrace{\sum_{t=1}^T \alpha_t f_t(x)}_{g(x)} \right)$$

Training Data Error Rate =

$$\begin{aligned} &= \frac{1}{N} \sum_n \delta(H(x^n) \neq \hat{y}^n) \\ &\leq \frac{1}{N} \sum_n \exp(-\hat{y}^n g(x^n)) \\ &= \prod_{t=1}^T 2\sqrt{\epsilon_t(1 - \epsilon_t)} \end{aligned}$$

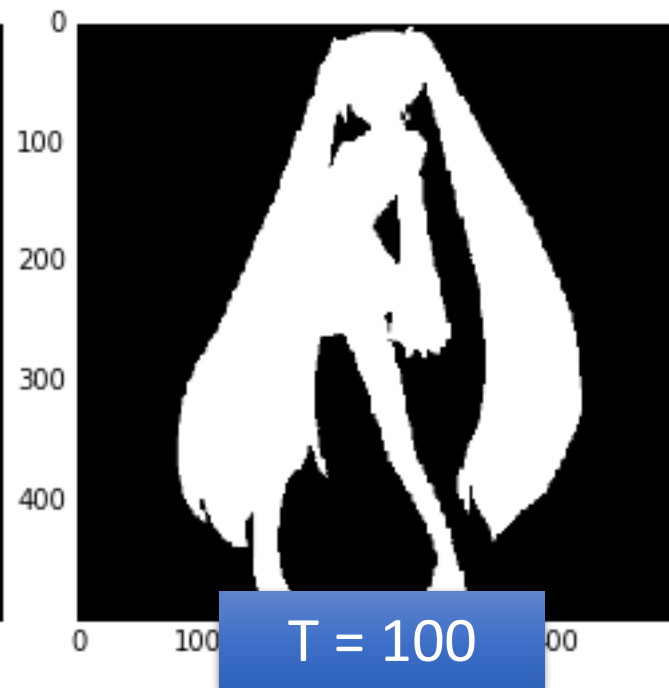
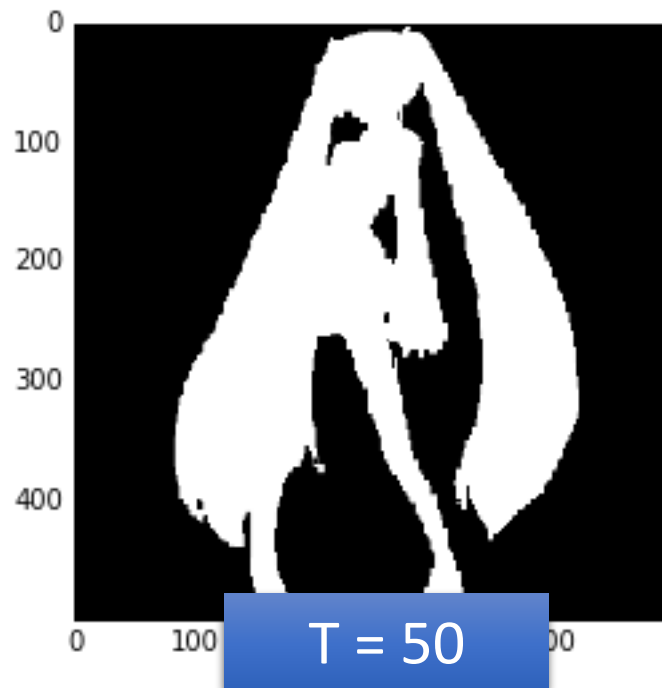
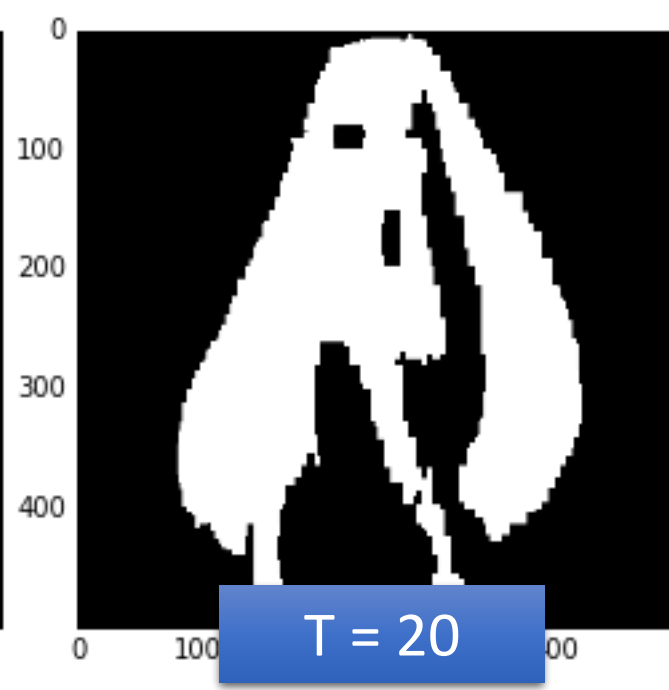
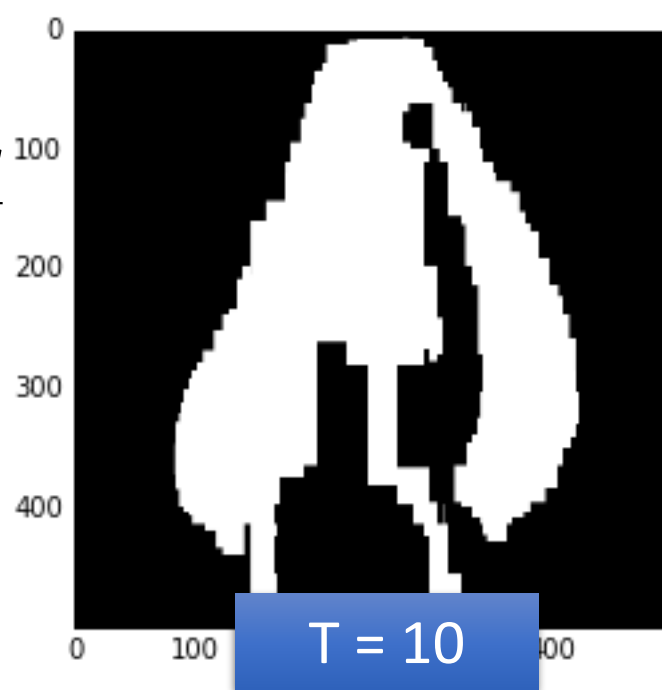
Getting smaller and smaller as T increase



**Experiment:**  
**Function of Miku**

Adaboost  
+Decision Tree

(depth = 5)



# To learn more ...

- Introduction of Adaboost:
  - Freund; Schapire (1999). "A Short Introduction to Boosting"
- Multiclass/Regression
  - Y. Freund, R. Schapire, "A Decision-Theoretic Generalization of on-Line Learning and an Application to Boosting", 1995.
  - Robert E. Schapire and Yoram Singer. Improved boosting algorithms using confidence-rated predictions. In Proceedings of the Eleventh Annual Conference on Computational Learning Theory, pages 80–91, 1998.
- Gentle Boost
  - Schapire, Robert; Singer, Yoram (1999). "Improved Boosting Algorithms Using Confidence-rated Predictions".

# General Formulation of Boosting

- Initial function  $g_0(x) = 0$
- For  $t = 1$  to  $T$ :
  - Find a function  $f_t(x)$  and  $\alpha_t$  to improve  $g_{t-1}(x)$ 
    - $g_{t-1}(x) = \sum_{i=1}^{t-1} \alpha_i f_i(x)$
    - $g_t(x) = g_{t-1}(x) + \alpha_t f_t(x)$
- Output:  $H(x) = \text{sign}(g_T(x))$

What is the learning target of  $g(x)$ ?

$$\text{Minimize } L(g) = \sum_n l(\hat{y}^n, g(x^n)) = \sum_n \exp(-\hat{y}^n g(x^n))$$

# Gradient Boosting

- Find  $g(x)$ , minimize  $L(g) = \sum_n \exp(-\hat{y}^n g(x^n))$ 
  - If we already have  $g(x) = g_{t-1}(x)$ , how to update  $g(x)$ ?

Gradient Descent:

$$g_t(x) = g_{t-1}(x) - \eta \left. \frac{\partial L(g)}{\partial g(x)} \right|_{g(x) = g_{t-1}(x)}$$

$\sum_n \exp(-\hat{y}^n g_{t-1}(x^n)) (-\hat{y}^n)$

Same direction

$$g_t(x) = g_{t-1}(x) + \alpha_t f_t(x)$$

# Gradient Boosting

$$f_t(x) \xleftrightarrow{\text{Same direction}} \sum_n \exp(-\hat{y}^n g_t(x^n)) (\hat{y}^n)$$

We want to find  $f_t(x)$  maximizing

$$\sum_n \underbrace{\exp(-\hat{y}^n g_{t-1}(x^n))}_{\text{example weight } u_t^n} \underbrace{(\hat{y}^n) f_t(x^n)}_{\text{Minimize Error, Same sign}}$$

$$u_t^n = \exp(-\hat{y}^n g_{t-1}(x^n)) = \exp\left(-\hat{y}^n \sum_{i=1}^{t-1} \alpha_i f_i(x^n)\right)$$

$$= \prod_{i=1}^{t-1} \exp(-\hat{y}^n \alpha_i f_i(x^n))$$

Exactly the weights we obtain in Adaboost

# Gradient Boosting

- Find  $g(x)$ , minimize  $L(g) = \sum_n \exp(-\hat{y}^n g(x^n))$

$$g_t(x) = g_{t-1}(x) + \alpha_t f_t(x)$$

$\alpha_t$  is something like  
learning rate

Find  $\alpha_t$  minimizing  $L(g_{t+1})$

$$L(g) = \sum_n \exp(-\hat{y}^n (g_{t-1}(x) + \alpha_t f_t(x)))$$

$$= \sum_n \exp(-\hat{y}^n g_{t-1}(x)) \exp(-\hat{y}^n \alpha_t f_t(x))$$

$$= \sum_{\hat{y}^n \neq f_t(x)} \exp(-\hat{y}^n g_{t-1}(x^n)) \exp(\alpha_t) + \sum_{\hat{y}^n = f_t(x)} \exp(-\hat{y}^n g_{t-1}(x^n)) \exp(-\alpha_t)$$

Find  $\alpha_t$   
such that

$$\frac{\partial L(g)}{\partial \alpha_t} = 0$$

$$\alpha_t = \ln \sqrt{(1 - \varepsilon_t) / \varepsilon_t}$$

Adaboost!

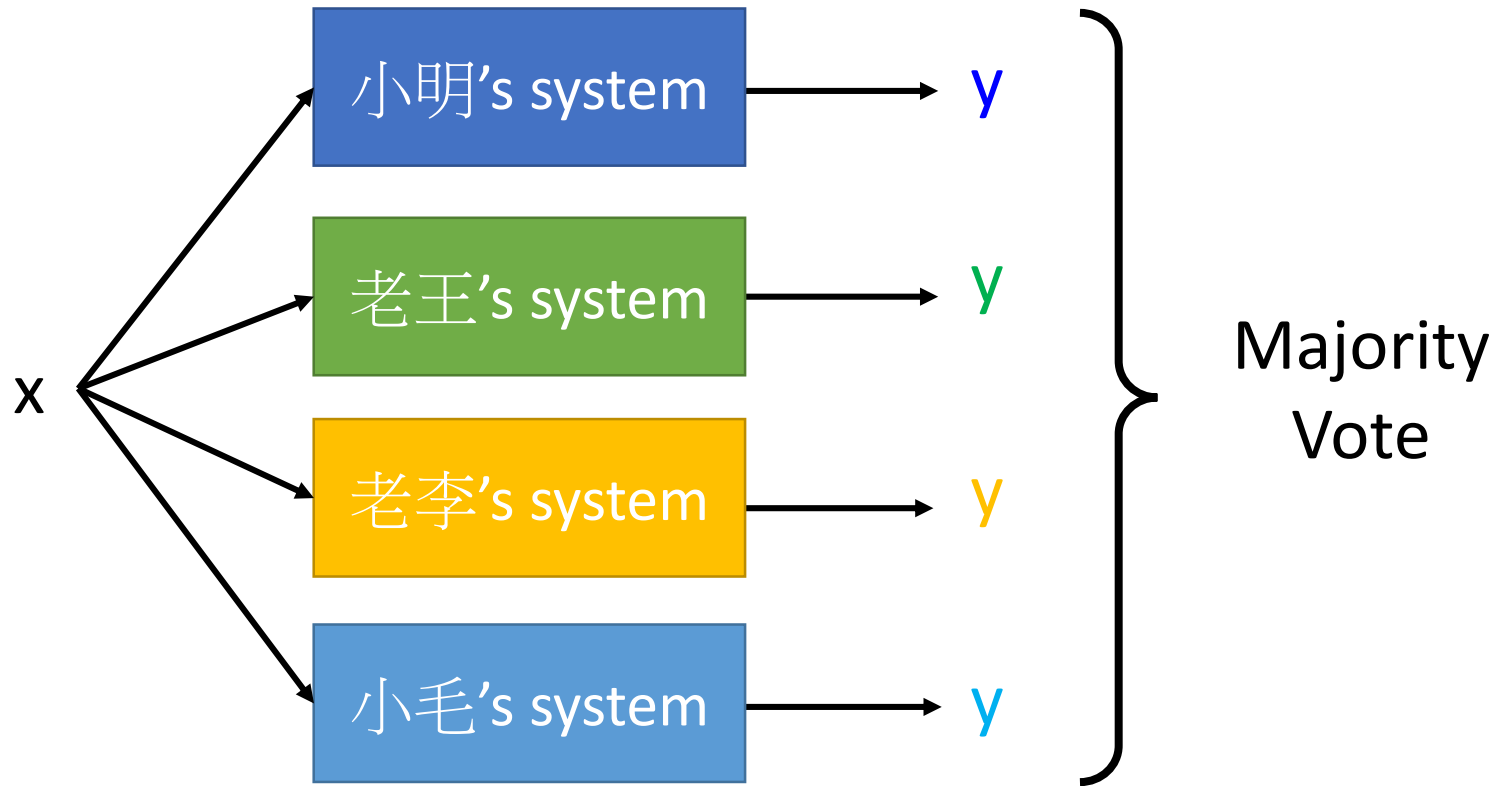
# Cool Demo

- [http://arogozhnikov.github.io/2016/07/05/gradient\\_boosting\\_playground.html](http://arogozhnikov.github.io/2016/07/05/gradient_boosting_playground.html)



# Ensemble: Stacking

# Voting



# Stacking

